

WEEK 2

TIDYVERSE & MANIPULATING DATA

DATA VISUALIZATION FOR SOCIAL SCIENTISTS

LECTURER: JEFFREY ZIEGLER, PhD

TEACHING FELLOW: SHEKHAR KEDIA

ASDS - TRINITY COLLEGE DUBLIN

SPRING 2026

ROAD MAP FOR TODAY

■ Today:

- ▶ Work with rectangular data using gapminder dataset
- ▶ Manipulate data with core dplyr verbs
- ▶ Summarize amounts and proportions and plot them with ggplot2
 - Practice histograms, density plots, ridge plots, boxplots, violins, and jittered points
 - Emphasize choices (bin width, bandwidth, outlier rules) that change how uncertainty appears

■ By next week, please...

- ▶ Problem set #2

KEY DPLYR VERBS

- **filter()**: extract rows (cases) that meet logical conditions
- **select()**: choose columns (variables) from a data frame
- **arrange()**: sort rows by one or more variables
- **mutate()**: create or transform columns
- **group_by() + summarize()**: compute group-level summaries

FILTER(): SUBSETTING ROWS

- Pattern: `filter(data, condition)` returns rows where condition is TRUE
- Example: `filter(gapminder, country == "Denmark")` returns all Denmark rows
- Uses logical operators such as `==, <, >, <=, >=`

LOGICAL TESTS

- $x < y$: less than
- $x > y$: greater than
- $x == y$: equal to (used for testing, not assignment)
- $x \leq y, x \geq y$: less/greater than or equal to

EX: STUDYING AFRICAN FARMER-LED IRRIGATION

- SAIFI study looked at farming and irrigation methods in Tanzania and Mozambique
- Survey data was collected through interviews conducted between November 2016 and June 2017

Variable	Description
village	Village name
interview_date	Date of interview
no_membrs	How many members in the household?
years_liv	How many years have you been living in this village or neighboring village?
rooms	How many rooms in the main house are used for sleeping?
memb_assoc	Are you a member of an irrigation association?
affect_conflicts	Have you been affected by conflicts with other irrigators in the area?
liv_count	Number of livestock owned
no_meals	How many meals do people in your household normally eat in a day?

Ex: FILTERING WITH FILTER()

```
1 # Import dataset and load library
2 SAIFI <- read.csv("https://raw.githubusercontent.com/ASDS-TCD/
  DataViz_2026/refs/heads/main/datasets/SAIFI.csv")
3
4 # All households in village "God"
5 SAIFI_god <- filter(SAIFI, village == "God")
6
7 # Households with more than 6 members
8 large_households <- filter(SAIFI, no_membrs > 6)
9
10 # Households in "God" or "Ruaca" with more than 4 members
11 god_ruaca_large <- filter(SAIFI, village %in% c("God", "Ruaca") &
  no_membrs > 4)
```

Ex: SELECTING WITH SELECT()

```
1 # Keep only identification and location variables
2 id_loc <- SAIFI |>
3   select(key_ID, village)
4
5 # Move village to the first position
6 village_first <- SAIFI |>
7   select(village, everything())
8
9 # Drop columns related to survey timing
10 no_timing <- SAIFI |>
11   select(-starts_with("interview_date"))
```

Ex: ARRANGING DATA WITH ARRANGE()

```
1 # Households ordered by size (ascending)
2 SAIFI_by_size <- SAIFI |>
3   arrange(no_membrs)
4
5 # Households ordered by size (descending) within village
6 SAIFI_by_village_size <- SAIFI |>
7   arrange(village , desc(no_membrs))
```

MUTATE(): CREATING NEW VARIABLES

- Pattern: `mutate(data, new_var = expression)` adds or modifies columns
- Examples: indicator for African countries, logged GDP per capita, classification of regions
- Often combined with `group_by()` and `summarize()` for richer summaries

Ex: ADDING NEW COLUMNS WITH MUTATE()

```
1 # Add a logical indicator for "large" households
2 SAIFI <- SAIFI |>
3   mutate(large_hh = no_membrs >= 6)
4
5 # Years of living in village in decades
6 SAIFI <- SAIFI |>
7   mutate(living_decades = years_liv / 10)
```

Ex: COMBINING MULTIPLE VERBS WITH PIPES (|>)

```
1 # Categorize household size
2 SAFI <- SAFI |>
3   mutate(
4     hh_size_cat = case_when(
5       no_membrs <= 3 ~
6         "small",
7       no_membrs >= 4 & no_membrs ~
8         <= 7 ~ "medium",
9       no_membrs > 7 ~
10        "large"
11      )
12    )
```

```
1 # Large households in Ruaca with
2   selected columns
3 ruaca_large_small_df <- SAFI |>
4   filter(village == "Ruaca", no_
5     membrs > 6) |>
6   mutate(
7     large_hh = TRUE,
8     members_per_room = no_membrs /
9       rooms
10    ) |>
11   select(village, no_membrs,
12     rooms, members_per_room,
13     large_hh)
```

TWO-STEP WORKFLOW FOR PROPORTIONS

- Step 1: Summarize data with dplyr, often using group_by() and summarize()
- Step 2: Visualize summarized data with ggplot2
- Clarifies how proportions are computed before plotting them

SUMMARIZING WITH GROUP_BY() AND SUMMARIZE()

- Workflow: `data |> group_by(group_vars) |> summarize(stats)`
- Typical statistics: mean, median, min, max, count, etc.
- Used to compute quantities like median house members by village


```
1 # Mean household size by village
2 village_hh_summary <- SAFI |>
3   group_by(village) |>
4   summarize(
5     mean_members = mean(no_membrs, na.rm = TRUE),
6     median_members = median(no_membrs, na.rm = TRUE),
7     n_households = n(),
8     .groups = "drop"
9   )
```

Ex: PROPORTIONS WITH GROUP_BY() |> SUMMARIZE()

```
1 # Mean years in village by roof type
2 wall_living_summary <- SAIFI |>
3   count(wall_type, village, name = "n") |>
4   group_by(village) |>
5   mutate(prop = n / sum(n))
```

Ex: PROPORTIONS WITH GROUP_BY() |> SUMMARIZE()

```
1 ggplot(wall_living_summary,
2         aes(x = village , y =
3             prop , fill = wall_type)
4     ) +
5   geom_col(position = "fill")
6   +
7   scale_y_continuous(labels =
8       scales::percent_format
9   ()) +
10  labs(
11    x = "\nVillage",
12    y = "Proportion of
13        households\n",
14    fill = "Wall type",
15    title = "Wall material
16        proportions by village"
17  )
```



FROM SUMMARIES TO GRAPHICS


- Use summarized tables (e.g., counts or proportions by group) as ggplot inputs
- Choose appropriate geometries (bars, lines, etc.) to represent amounts and proportions
- Emphasize truthful, clear visual communication of aggregated data

HISTOGRAMS AND BIN WIDTH

- Histograms show counts of observations within ranges (bins) of a variable such as no_members
- Bin width strongly affects the apparent shape of the distribution and must be chosen deliberately
- Options include binwidth, bar border color, fill color (e.g. #E16462), and boundary to align bins at whole numbers


Ex: HISTOGRAM OF HOUSEHOLD SIZE

Default bin widths (=30)



```
1 ggplot(SAFI, aes(x = no_membrs)) +  
2   geom_histogram(color = "white", fill = "#  
3     E16462") +  
4   labs(x = "\nNumber of household members",  
      y = "Count\n")
```

bin=1




```
1 ggplot(SAFI, aes(x = no_membrs)) +  
2   geom_histogram(binwidth = 1, color = "  
3     white", fill = "#E16462") +  
4   labs(x = "\nNumber of household members",  
      y = "Count\n")
```

HISTOGRAMS WITH GROUPS


- Additional variables like `memb_assoc` (member of irrigation association) can be mapped to fill
- Overlaid histograms with multiple fills can be hard to interpret when groups overlap
- Faceting (e.g. `facet_wrap(vars(village))`) separates groups into small multiples for clearer comparisons

Ex: HISTOGRAMS WITH GROUPS

Member association



Member association by village



```
1 ggplot(SAFI, aes(x = no_membrs, fill = memb_
  assoc)) +
  2   geom_histogram(position="dodge") +
  3   labs(x = "\nNumber of household members",
  4     y = "Count\n")
```


```
1 ggplot(SAFI, aes(x = no_membrs, fill = memb_
  assoc)) +
  2   geom_histogram(position="dodge") +
  3   facet_wrap(vars(village)) +
  4   labs(x = "\nNumber of household members",
  5     y = "Count\n")
```

DENSITY AND RIDGE PLOTS

- Density plots provide a smoothed estimate of the distribution; bandwidth (bw) controls smoothing
- Aesthetics include fill (e.g. #E16462), border color (e.g. #9C3836), and line width
- Mapping fill = drv and using transparency (alpha) overlays group densities; ggridges can convert these to ridge plots with geom_density_ridges()


Ex: # OF HOUSEHOLD MEMBERS BY VILLAGE

Density plots



```
1 ggplot(SAFI, aes(x = no_membrs, fill =  
2   village)) +  
3   geom_density(alpha = 0.5, color = NA,  
4   adjust = 1) +  
5   labs(x = "\nNumber of household members",  
6       y = "Density\n",  
7       fill = "Village")
```

Ridge plots




```
1 ggplot(SAFI, aes(x = no_membrs, y = village ,  
2   fill = village)) +  
3   geom_density_ridges(alpha = 0.7, color = "  
4   white", scale = 1.2) +  
5   labs(x = "\nNumber of household members",  
6       y = "Village\n",  
7       fill = "Village")
```

BOXPLOTS, VIOLINS, AND DOTS

- Boxplots show quartiles, median, and outliers; coef controls the outlier rule (e.g. $5 \times \text{IQR}$)
- Visual settings include fill (e.g. #E6AD3C) and color (e.g. #5ABD51)
- Violin plots show mirrored density shapes; adding jittered points (e.g. width 0.1, size 0.5) reveals individual observations by category such as village


Ex: # OF HOUSEHOLD MEMBERS BY VILLAGE

Box plot



```
1 ggplot(SAFI, aes(x = village , y = no_membrs,
                  fill = village)) +
2   geom_boxplot(alpha = 0.7, color = "#5ABD51",
                  coef = 3, width = 0.4) +
3   labs(fill = "Village", x = "\nVillage", y
                  = "Household members\n")
```

Violin plot



```
1 ggplot(SAFI, aes(x = village , y = no_membrs,
                  fill = village)) +
2   geom_violin(alpha = 0.7, width = 0.8, draw
                  _quantiles = c(0.25, 0.5, 0.75)) +
3   geom_jitter(width = 0.15, alpha = 0.6,
                  size = 0.8) +
4   labs(fill = "Village", x = "\nVillage", y
                  = "Household members\n")
```

CLASS BUSINESS

- Read required (and suggested) online materials
- Fork GitHub repository
- Problem set # 1 is up on GitHub